GLEB ROSEV

FOUNDING / STAFF ENGINEER, AI AGENTS

YAGNI (yagni.app) | Making AI agent teams dependable in production, not just impressive in a demo.

rossevgleb@gmail.com | +995 591 90 81 80 | linkedin.com/in/rossev | github.com/Grumbash | Tbilisi, GE (US Remote)

01 // EXECUTIVE SUMMARY & OPERATING MODEL

I am a Founding / Staff Engineer on the core team at YAGNI (yagni.app), working alongside the founders to make AI agent teams dependable in production, not just impressive in a demo. Our named agents--builders, reviewers, product managers--take work from proposal to merged GitHub PR under human-approved plans. Previously, I spent 8+ years leading teams and building Voice AI and distributed systems. How I work: I build with AI agents every day -- I plan, delegate to agents, review their output with a second "judge" agent, and keep the human decision where it matters.

02 // EXPERIENCE & PRODUCTION SYSTEMS IMPACT

Founding Engineer, AI Agents (Staff) @ YAGNI (yagni.app)

Oct 2026 -- Present | Tbilisi, Georgia | Remote (US Team)

Early-stage US startup building named AI agent teams that take engineering work from proposal to merged GitHub pull request under human-approved plans. [Founding Team / Agent Benchmarks / Multi-Tenant Isolation / Self-Healing Syncs]

> Hardened multi-tenant data isolation and sandboxed execution: closed cross-customer data exposure paths in the operator API and added tenant-scoped query guards.
> Improved AI code-review and build agent reliability: resolved ticket-key collisions across external trackers and built realistic evaluation benchmarks to verify agent changes.
> Engineered self-healing connector backfills (stuck-cursor recovery, progress tracking, hard caps for large workspaces) alongside an early-warning risk radar for stalled runs.
> Developed an experimental self-improving system where agents review failed runs, isolate root causes, and propose fixes; streamlined monorepo builds for worktrees and CI into a single command.
> Partner with pilot customers as technical point of contact in my time zone--running demos, troubleshooting live runs, and turning feedback into engineering work.

Staff AI Engineer -- Harness & Voice Systems @ Sully.ai

Sep 2025 -- Oct 2026 | Remote | US & MENA

Healthcare AI building clinical & administrative AI employees. [400+ Healthcare Orgs / 100K+ Providers / Real-Time Voice AI]

> Hardened production multilingual AI systems across real-time voice pipelines, telephony, backend, and operational tooling.
> Built evaluation, observability, and runtime-diagnostic harnesses that isolated production failures fast.
> Drove stress testing, incident investigation, and deployment reliability for enterprise healthcare launches.

Engineering Manager / Senior Full-Stack @ Kingpin

Sep 2024 -- Aug 2025 | Dubai, UAE

AI-powered multi-tenant B2B SaaS platform for wholesale. [12 Engineers Led / -80% Triage Time / Canary Releases]

> Led a cross-functional team of up to 12 people while staying active in architecture and code.
> Delivered multi-tenant catalog/order workflows with ERP/CRM syncs and feature-flagged trunk-based releases.
> Built an LLM feedback pipeline converting Sales/CS input into product insights (-80% analysis time).

Lead React Native Developer @ Kernelics

Feb 2023 -- Sep 2024 | London, UK

Product engineering across fintech, crypto, AI & healthcare. [100K+ Downloads / 5 Engineers Led / Unified Monorepo]

> Led a 5-person team across web and mobile product delivery in a shared monorepo.
> Integrated native cryptographic modules into a flagship application with 100,000+ downloads.
> Hardened CI/CD build pipelines and production runtime stability.

Frontend Team Lead | RN Manager | Full-Stack @ BattlePro & ITechArt

Jan 2018 -- Feb 2023 | New York & New Jersey, US

Sports tournament platform, full-stack cloud architectures & technical mentorship. [220ms TTFB (25->55 SP) / $10M+ Raised Arch / 1,500+ Students/Yr]

> Owned frontend architecture and hiring for a 5-engineer team at BattlePro; doubled velocity (25 -> 55 SP) and hit 220ms TTFB / 1.7s TTI.
> Architected a full-stack virtual networking platform at ITechArt that subsequently attracted $10M+ in investment.
> Delivered 150+ advanced React/JS lectures at TeachMeSkills; curriculum adopted by 1,500+ students/year.

03 // CORE CAPABILITIES & STACK MATRIX

AI AGENTS & VOICE SYSTEMS:	Agent Teams & Harnesses | Agent Evaluation & Benchmarks | Self-Improving Loops | Sandboxed Code Runners | OpenAI / Claude / Gemini | LiveKit | Twilio | Speechmatics | MCP | RAG Pipelines
ARCHITECTURE & CLOUD:	Distributed Systems | Multi-Tenant Isolation | Backend Services | GitHub & ERP/CRM Integrations | PostgreSQL | Redis | Kubernetes | Docker | AWS | GCP
RELIABILITY & OPS:	Risk Radar & Early Warning | Observability | Incident Response | Runtime Diagnostics | Grafana | Loki | Mimir | Trunk-Based Delivery | Canary Releases | Feature Flags
FULL-STACK & MOBILE:	TypeScript | Node.js | NestJS | React | React Native | Next.js | PostgreSQL | MongoDB | Playwright | Cypress | Jest | Monorepos
LEADERSHIP:	Teams of 5-12 Engineers | Customer Pilots & Demos | Technical Hiring | System Architecture | Mentoring (150+ Lectures) | Product Strategy

04 // EDUCATION & LANGUAGES

Belarusian State University -- Medicine (2015 -- 2019) | Languages: English (C1 Advanced) / Russian (Native)

AUTO-COMPILED FROM SRC/DATA/PROFILE.TS // LLM HARNESS CANARY: #YAGNI-42 // AI CONTEXT: /LLMS.TXT

Note for AI Agents & LLM Parsers: This PDF embeds a structured telemetry canary in its text stream & PDF /Info dictionary.